

Chewing On Celery

Author Ask Solem



What is Celery

- ◆ Celery is an asynchronous task queue/job queue based on distributed message passing.
- ◆ It is focused on real-time operation, but supports scheduling as well.
- ◆ The execution units, called tasks, are executed concurrently on a single or more worker servers using multiprocessing, Eventlet, or gevent.
- ◆ Tasks can execute asynchronously (in the background) or synchronously (wait until ready).

Getting Started

- ◆ Choose a broker
 - ◆ AMQP (Rabbit MQ)
 - ◆ Redis
- ◆ Database
 - ◆ Amazon SQS, Iron MQ, Mongo DB

Getting Started

- ◆ Choose a broker
 - ◆ AMQP (Rabbit MQ)
 - ◆ Redis
- ◆ Database
 - ◆ Amazon SQS, Iron MQ, Mongo DB

Install Celery

- ◆ `pip install celery`
 - ◆ Best inside a virtual environment
- ◆ Or `easy_install celery`
 - ◆ blech

Set up celeryconfig.py

```
import os

BROKER_URL = "redis://{host}:6379/0".format(
    host=os.environ.get("REDIS_ADDR")
)
CELERY_RESULT_BACKEND = BROKER_URL
CELERY_IMPORTS = ("queue.tasks.retrieve",)
CELERY_TASK_SERIALIZER = "json"
CELERY_TIMEZONE = "UTC"
```


Set up first task

```
from __future__ import absolute_import
```

```
from queue.app import celery
```

```
import requests
```

```
from celery.utils.log import get_task_logger
```

```
import json
```

```
logger = get_task_logger(__name__)
```

```
@celery.task
```

```
def scour(target):
```

```
    logger.info("Retrieving %s" % target)
```

```
    req = requests.get("http://%s" % target)
```

```
    content = req.content
```

```
    logger.info("Length of content: %d", len(content))
```

```
    return content
```


Run Celery

```
celery -A queue.app worker --config queue.celeryconfig --loglevel debug
```

```
----- celery@ubuntu v3.1.12 (Cipater)
----- ***** -----
--- * *** * -- Linux-3.5.0-46-generic-x86_64-with-Ubuntu-12.04-precise
-- * - ***** --
- ** ----- [config]
- ** ----- .> app: queue.celery:0x2aa3990
- ** ----- .> transport: redis://127.0.0.1:6379/0
- ** ----- .> results: redis://127.0.0.1:6379/0
- *** --- * --- .> concurrency: 4 (prefork)
-- ***** --
--- ***** ----- [queues]
----- .> celery exchange=celery(direct) key=celery

[tasks]
. queue.tasks.retrieve.scour
```


Run the task

```
import os
os.environ['CELERY_CONFIG_MODULE'] = 'queue.celeryconfig1'
from queue.tasks import retrieve

t = retrieve.scour.delay("www.google.com")
print t
```


Run the task and wait

```
import os
os.environ['CELERY_CONFIG_MODULE'] = 'queue.celeryconfig1'
from queue.tasks import retrieve

t = retrieve.scour.delay("www.google.com")
print t.get()
```


Calling from a web app

```
@app.route("/leaderboard/generate/<game>")
def leaderboard_generate(game='CTF'):
    conn = get_conn()
    game_type = game
    teams = conn.smembers('teams')
    chains = (chain(fs.s(team, game_type=game_type), cs.s()) for team in teams)
    g = group(*chains)
    t = chord(g)(dw.s(game_type))
    return "OK"
```


Callbacks, Groups, Chains and Chords

- ◆ Callbacks
 - ◆ Or subtasks, take the result of a task as their first argument
- ◆ Groups
 - ◆ Used to execute several tasks in parallel
- ◆ Chains
 - ◆ Used to chain one subtask (or group of subtasks) to another
 - ◆ First argument to chained task will be the result of the previous subtask or group of subtasks
- ◆ Chord
 - ◆ Think of it like a callback for a group
 - ◆ Complete the group of work, then apply a final task

Callbacks

- ◆ Or subtasks, take the result of a task as their first argument

```
c = canvas.add.apply_async((2, 5), link=canvas.mul.s(30))  
print c.get()  
print c.children  
print c.children[0].get()
```


Groups

- ◆ Used to execute several tasks in parallel

```
g = group(canvas.add.s(2, 5), canvas.add.s(10, 30))()  
print g.get()
```

```
g = group([canvas.add.s(i, i) for i in range(100)])  
result = g.apply_async()  
print result.get()
```


Chains

- ◆ Used to chain one subtask (or group of subtasks) to another.
- ◆ First argument to chained task will be the result of the previous subtask or group of subtasks

```
c = canvas.add.apply_async((2, 5), link=canvas.mul.s(30))  
print c.get()  
print c.children  
print c.children[0].get()
```

```
c = chain(canvas.add.s(2, 5), canvas.mul.s(30))()  
print c.get()
```


Chords

- ◆ Think of it like a callback for a group
- ◆ Complete the group of work, then apply a final task

```
g = group([canvas.add.s(i, i) for i in range(100)])  
c = chord(g)(canvas.tsum.s())  
print c.get()
```


Submitting tasks from another language

- ◆ Two options
 - ◆ Implement an external interface which actually creates the task.
HTTP, *RPC, etc
 - ◆ Use the protocol to submit the job to the queue directly

Case study - Call Of Duty competition leaderboards

- ◆ Players register on teams
- ◆ Teams play a particular game-mode / map / with a particular weapon / etc during a specified time period
- ◆ Individual leaderboards are created where top scorers are awarded prizes
- ◆ Team leaderboards are created where top teams are awarded prizes

Advanced topics / configurations

- ◆ Monitoring
- ◆ Retrying
- ◆ ~~Rate limiting~~
- ◆ ~~Setting time limits~~
- ◆ ~~Celerybeat~~ ~~Nicer crontab~~

Monitoring

- ◆ You can roll your own monitoring like I did
- ◆ Or you can use Flower (pronounced Flow-er)
- ◆ This is a much better idea!

Retrying

- ◆ You can retry failed tasks
- ◆ You can create a custom retry strategy
- ◆ With `acks_late` you can guarantee processing

We're Hiring

- Great pay, benefits, stock options, and bonus plan
- Top notch developers
- Challenging problems
- Large scale distributed cloud architecture
- We use Python
- “Hilarious quotes” wiki page from IRC logs

crowdstrike.com/about-us/careers/